

# cnn

March 8, 2024

## 0.1 Cancer Detection using CNN's

I will use a CNN model to identify metastatic cancer in small image patches taken from larger digital pathology scans. The data for this competition is a slightly modified version of the Patch-Camelyon (PCam) benchmark dataset (the original PCam dataset contains duplicate images due to its probabilistic sampling, however, the version presented on Kaggle does not contain duplicates). The data has 440,050 images in the training set and 114,916 images in validation (test) set.

### 0.1.1 EDA

The first step of the EDA will create data structures using pandas.

```
[ ]: import pandas as pd
import cv2
import matplotlib.pyplot as plt
import os
import numpy as np

[ ]: df_data = pd.read_csv('data/histopathologic-cancer-detection/train_labels.csv')

print(len(os.listdir('data/histopathologic-cancer-detection/train')))
print(len(os.listdir('data/histopathologic-cancer-detection/test')))
```

440050  
57458

The next step is to visualize the some of the images from both categories so that we can visually see examples of the images we are trying to train.

```
[ ]: # source: https://www.kaggle.com/gpreda/honey-bee-subspecies-classification

def draw_category_images(col_name, figure_cols, df, IMAGE_PATH):

    """
    Give a column in a dataframe,
    this function takes a sample of each class and displays that
    sample on one row. The sample size is the same as figure_cols which
    is the number of columns in the figure.
```

Because this function takes a random sample, each time the function is run, it

displays different images.

"""

```
df[col_name]= df[col_name].astype(float)
categories = (df.groupby([col_name])[col_name].nunique()).index
f, ax = plt.subplots(nrows=len(categories),ncols=figure_cols,
                    figsize=(4*figure_cols,4*len(categories))) # adjust
```

size here

# draw a number of images for each location

```
for i, cat in enumerate(categories):
```

```
    sample = df[df[col_name]==cat].sample(figure_cols) # figure_cols is
```

also the sample size

```
    for j in range(0,figure_cols):
```

```
        file=IMAGE_PATH + sample.iloc[j]['id'] + '.tif'
```

```
        im=cv2.imread(file)
```


```
        ax[i, j].imshow(im, resample=True, cmap='gray')
```

```
        ax[i, j].set_title(cat, fontsize=16)
```

```
plt.tight_layout()
```

```
plt.show()
```

```
[ ]: draw_category_images('label',4, df_data,'data/histopathologic-cancer-detection/
train/')
```



Then create a training and testing splits of the data in preparation for training

```
[ ]: from sklearn.model_selection import train_test_split

y = df_data['label']
df_train, df_val = train_test_split(df_data, test_size=0.10, random_state=101,
    ↪stratify=y)

print(df_train)
print(df_val)
print(y)
```

|        | id                                       | label |
|--------|--|-------|
| 134709 | 2ddb582347063e13f6e8e8bcc693c2ddc65f5781 | 0.0   |
| 122063 | 9563cbc9f43d2d2ecc1e8b757ee85637db430fe7 | 0.0   |
| 171972 | 2ca2f3a7c24ba5c5638089e1c08088ac246454fe | 1.0   |
| 63378  | ab452f00d0fc2a3dd7c4781c0dcaaecf4d71e4ab | 1.0   |
| 36782  | 85aaec3a61c47e33d4bad167f9c1a013a2ff2401 | 0.0   |
| ...    | ...                                      | ...   |
| 117398 | 9e04c5701b81604b04bd340fb86f26c0ec0fa84d | 0.0   |
| 87646  | 86611a3dae286bf57808cce6bd952777a83e2f98 | 0.0   |
| 99059  | 7344089a9732d851d3a5d41cf9c7849e53fc1e94 | 0.0   |
| 121588 | a49c538708c85c9be22946c40e284d357bf960b5 | 0.0   |
| 21829  | d7a3e35ffb2dfa20db4aff2c9e91edb64f4a369c | 0.0   |

[198022 rows x 2 columns]

|        | id                                       | label |
|--------|--|-------|
| 79961  | 0d8c9abf388785a016d72e609ab932b7e134f827 | 0.0   |
| 45724  | 5566aa9c1d698535c449144d72e98556bf77fd47 | 0.0   |
| 81384  | 358c0ab1962f8f1c202b146d88b35f55252bde9f | 1.0   |
| 91989  | 00f81cefedcc1fb79a5c500677aad14438c50296 | 0.0   |
| 177529 | e7a806a270062e1bf77aed9be266953d11e25bd6 | 0.0   |
| ...    | ...                                      | ...   |
| 190249 | f88d7fe078e1b12862a2b96965b01045652483ff | 0.0   |
| 105101 | 6cf9e9b14069ebcc195ad4cfbab1f3d623be403f | 1.0   |
| 19531  | dcc0d91a6022f62cecdf37c1e1978f168eb36ae7 | 0.0   |
| 119918 | b9174d52655f0c0a712cfe81aac2f12cde46c165 | 0.0   |
| 215077 | f407c348a2d5d29f670c9fe58fc045e683fdef0e | 0.0   |

[22003 rows x 2 columns]

|        |     |
|--------|-----|
| 0      | 0.0 |
| 1      | 1.0 |
| 2      | 0.0 |
| 3      | 0.0 |
| 4      | 0.0 |
| ...    |     |
| 220020 | 0.0 |
| 220021 | 1.0 |
| 220022 | 0.0 |

```
220023    0.0
220024    1.0
Name: label, Length: 220025, dtype: float64
```

```
[ ]: import os
import shutil

base_dir = 'data/histopathologic-cancer-detection/split'
os.mkdir(base_dir)

train_dir = os.path.join(base_dir, 'train_dir')
os.mkdir(train_dir)

val_dir = os.path.join(base_dir, 'val_dir')
os.mkdir(val_dir)

no_tumor_tissue = os.path.join(train_dir, 'tumor_0')
os.mkdir(no_tumor_tissue)
has_tumor_tissue = os.path.join(train_dir, 'tumor_1')
os.mkdir(has_tumor_tissue)

# create new folders inside val_dir
no_tumor_tissue = os.path.join(val_dir, 'tumor_0')
os.mkdir(no_tumor_tissue)
has_tumor_tissue = os.path.join(val_dir, 'tumor_1')
os.mkdir(has_tumor_tissue)

df_data_ind = df_data.set_index('id')

train_list = list(df_train['id'])
val_list = list(df_val['id'])

for image in train_list:

    # the id in the csv file does not have the .tif extension therefore we add
    ↳ it here
    fname = image + '.tif'
    # get the label for a certain image
    target = df_data_ind.loc[image, 'label']

    # these must match the folder names
    if target == 0:
        label = 'tumor_0'
    if target == 1:
        label = 'tumor_1'
```

```

# source path to image
src = os.path.join('data/histopathologic-cancer-detection/train', fname)
# destination path to image
dst = os.path.join(train_dir, label, fname)
# copy the image from the source to the destination
shutil.copyfile(src, dst)

# Transfer the val images

for image in val_list:

    # the id in the csv file does not have the .tif extension therefore we add_
    ↪ it here
    fname = image + '.tif'
    # get the label for a certain image
    target = df_data_ind.loc[image, 'label']

    # these must match the folder names
    if target == 0:
        label = 'tumor_0'
    if target == 1:
        label = 'tumor_1'

    # source path to image
    src = os.path.join('data/histopathologic-cancer-detection/train', fname)
    # destination path to image
    dst = os.path.join(val_dir, label, fname)
    # copy the image from the source to the destination
    shutil.copyfile(src, dst)

```

### 0.1.2 Model Architecture

The first step of the model architecture is to use tensorflow's ImageDataGenerator tool to structure the data into files for easier modeling using and training.

```

[ ]: from keras.preprocessing.image import ImageDataGenerator

train_path = 'data/histopathologic-cancer-detection/split/train_dir'
valid_path = 'data/histopathologic-cancer-detection/split/val_dir'
test_path = 'data/histopathologic-cancer-detection/test'

img_size = 96
num_train_samples = len(df_train)
num_val_samples = len(df_val)
train_batch_size = 32

```

```

val_batch_size = 32

train_steps = np.ceil(num_train_samples / train_batch_size)
val_steps = np.ceil(num_val_samples / val_batch_size)

datagen = ImageDataGenerator(rescale=1.0/255)

datagen = ImageDataGenerator(preprocessing_function=lambda x:(x - x.mean()) / x.
    ↪std() if x.std() > 0 else x,
                             horizontal_flip=True,
                             vertical_flip=True)

train_gen = datagen.flow_from_directory(train_path,
                                       target_size=(img_size,img_size),
                                       batch_size=train_batch_size,
                                       class_mode='binary')

val_gen = datagen.flow_from_directory(valid_path,
                                     target_size=(img_size,img_size),
                                     batch_size=val_batch_size,
                                     class_mode='binary')

# Note: shuffle=False causes the test dataset to not be shuffled
test_gen = datagen.flow_from_directory(valid_path,
                                       target_size=(img_size,img_size),
                                       batch_size=1,
                                       class_mode='categorical',
                                       shuffle=False)

```

Found 198022 images belonging to 2 classes.

Found 22003 images belonging to 2 classes.

Found 22003 images belonging to 2 classes.

```

[ ]: import tensorflow as tf
from tensorflow.keras.layers import Conv2D, MaxPool2D,MaxPooling2D
from tensorflow.keras.layers import Dense, Dropout, Flatten,
    ↪BatchNormalization, Activation
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import EarlyStopping, ReduceLRonPlateau,
    ↪ModelCheckpoint
from tensorflow.keras.optimizers import Adam
from keras.callbacks import EarlyStopping, ReduceLRonPlateau

```

I will then create a few callback functions that will stop training after if the training plateaus.

```
[ ]: earlystopper = EarlyStopping(monitor='val_loss', patience=2, verbose=1,
    ↪ restore_best_weights=True)
reducel = ReduceLROnPlateau(monitor='val_loss', patience=1, verbose=1, factor=0.
    ↪ 1)

callbacks_list = [reducel, earlystopper]
```

```
[ ]: kernel_size = (3,3)
pool_size= (2,2)
first_filters = 32
second_filters = 64
third_filters = 128

dropout_conv = 0.3
dropout_dense = 0.5
```

Convolutional Neural Networks are the obvious choice for a problem of this type as we can use supervised learning and want to identify patterns in images that are not easily identified by human visuals.

To start I will create a CNN that relies on relu algorithms and filters that to help identify different features in the images. The dropout layers help limit overfitting by limiting tensors from co-adapting too much.

```
[ ]: model1 = Sequential()
model1.add(Conv2D(first_filters, kernel_size, activation = 'relu', input_shape=
    ↪ (96, 96, 3)))
model1.add(Conv2D(first_filters, kernel_size, activation = 'relu'))
model1.add(Conv2D(first_filters, kernel_size, activation = 'relu'))
model1.add(MaxPooling2D(pool_size = pool_size))
model1.add(Dropout(dropout_conv))

model1.add(Conv2D(second_filters, kernel_size, activation = 'relu'))
model1.add(Conv2D(second_filters, kernel_size, activation = 'relu'))
model1.add(Conv2D(second_filters, kernel_size, activation = 'relu'))
model1.add(MaxPooling2D(pool_size = pool_size))
model1.add(Dropout(dropout_conv))

model1.add(Conv2D(third_filters, kernel_size, activation = 'relu'))
model1.add(Conv2D(third_filters, kernel_size, activation = 'relu'))
model1.add(Conv2D(third_filters, kernel_size, activation = 'relu'))
model1.add(MaxPooling2D(pool_size = pool_size))
model1.add(Dropout(dropout_conv))

model1.add(Flatten())
model1.add(Dense(256, activation = "relu"))
model1.add(Dropout(dropout_dense))
```

```
model1.add(Dense(1, activation = "softmax"))

model1.summary()
```

Model: "sequential"

| Layer (type)                   | Output Shape        | Param # |
|--------------------------------|---------------------|---------|
| conv2d (Conv2D)                | (None, 94, 94, 32)  | 896     |
| conv2d_1 (Conv2D)              | (None, 92, 92, 32)  | 9248    |
| conv2d_2 (Conv2D)              | (None, 90, 90, 32)  | 9248    |
| max_pooling2d (MaxPooling2D)   | (None, 45, 45, 32)  | 0       |
| dropout (Dropout)              | (None, 45, 45, 32)  | 0       |
| conv2d_3 (Conv2D)              | (None, 43, 43, 64)  | 18496   |
| conv2d_4 (Conv2D)              | (None, 41, 41, 64)  | 36928   |
| conv2d_5 (Conv2D)              | (None, 39, 39, 64)  | 36928   |
| max_pooling2d_1 (MaxPooling2D) | (None, 19, 19, 64)  | 0       |
| dropout_1 (Dropout)            | (None, 19, 19, 64)  | 0       |
| conv2d_6 (Conv2D)              | (None, 17, 17, 128) | 73856   |
| conv2d_7 (Conv2D)              | (None, 15, 15, 128) | 147584  |
| conv2d_8 (Conv2D)              | (None, 13, 13, 128) | 147584  |
| max_pooling2d_2 (MaxPooling2D) | (None, 6, 6, 128)   | 0       |
| dropout_2 (Dropout)            | (None, 6, 6, 128)   | 0       |
| flatten (Flatten)              | (None, 4608)        | 0       |
| dense (Dense)                  | (None, 256)         | 1179904 |
| dropout_3 (Dropout)            | (None, 256)         | 0       |



dense\_1 (Dense) (None, 1) 257

```
=====
Total params: 1660929 (6.34 MB)
Trainable params: 1660929 (6.34 MB)
Non-trainable params: 0 (0.00 Byte)
-----
```

```
[ ]: model1.compile(Adam(learning_rate=0.0001), loss='binary_crossentropy',
                    metrics=['accuracy'])

filepath1 = "model.h5"

history1 = model1.fit(train_gen, steps_per_epoch=train_steps,
                      validation_data=val_gen,
                      validation_steps=val_steps,
                      epochs=10, verbose=1,
                      callbacks=callbacks_list)
```



```
[ ]: acc = history1.history['accuracy']
      val_acc = history1.history['val_accuracy']
      loss = history1.history['loss']
      val_loss = history1.history['val_loss']

      epochs = range(1, len(acc) + 1)

      plt.plot(epochs, acc, 'bo', label='Training Accuracy')
      plt.plot(epochs, val_acc, 'b', label='Validation Accuracy')
      plt.title('Training and Validation Accuracy')
      plt.legend()
      plt.figure()

      plt.plot(epochs, loss, 'bo', label='Training loss')
      plt.plot(epochs, val_loss, 'b', label='Validation loss')
      plt.title('Training and Validation loss')
      plt.legend()
      plt.figure()
```

```
[ ]: <Figure size 432x288 with 0 Axes>
```



<Figure size 432x288 with 0 Axes>

As we can see here while loss continues to go down the accuracy is very stagnate and has very poor

validation accuracy.

In the second model, we add batch normalization which helps eliminate bias by resetting the mean and standard deviations.

```
[ ]: model2 = Sequential()
model2.add(Conv2D(first_filters, kernel_size, activation = 'relu', input_shape=(img_size, img_size, 3)))
model2.add(Conv2D(first_filters, kernel_size, use_bias=False))
model2.add(BatchNormalization())
model2.add(Activation("relu"))
model2.add(MaxPool2D(pool_size = pool_size))
model2.add(Dropout(dropout_conv))

model2.add(Conv2D(second_filters, kernel_size, use_bias=False))
model2.add(BatchNormalization())
model2.add(Activation("relu"))
model2.add(Conv2D(second_filters, kernel_size, use_bias=False))
model2.add(BatchNormalization())
model2.add(Activation("relu"))
model2.add(MaxPool2D(pool_size = pool_size))
model2.add(Dropout(dropout_conv))

model2.add(Conv2D(third_filters, kernel_size, use_bias=False))
model2.add(BatchNormalization())
model2.add(Activation("relu"))
model2.add(Conv2D(third_filters, kernel_size, use_bias=False))
model2.add(BatchNormalization())
model2.add(Activation("relu"))
model2.add(MaxPool2D(pool_size = pool_size))
model2.add(Dropout(dropout_conv))

#model2.add(GlobalAveragePooling2D())
model2.add(Flatten())
model2.add(Dense(256, use_bias=False))
model2.add(BatchNormalization())
model2.add(Activation("relu"))
model2.add(Dropout(dropout_dense))
model2.add(Dense(1, activation = "sigmoid"))

model2.summary()
```

Model: "sequential\_2"

| Layer (type)      | Output Shape       | Param # |
|-------------------|--------------------|---------|
| conv2d_6 (Conv2D) | (None, 94, 94, 32) | 896     |

|   |                     |         |
|---|---------------------|---------|
| conv2d_7 (Conv2D)                           | (None, 92, 92, 32)  | 9216    |
| batch_normalization_6 (Batch Normalization) | (None, 92, 92, 32)  | 128     |
| -----                                       |                     |         |
| Layer (type)                                | Output Shape        | Param # |
| =====                                       |                     |         |
| conv2d_6 (Conv2D)                           | (None, 94, 94, 32)  | 896     |
| conv2d_7 (Conv2D)                           | (None, 92, 92, 32)  | 9216    |
| batch_normalization_6 (Batch Normalization) | (None, 92, 92, 32)  | 128     |
| activation_6 (Activation)                   | (None, 92, 92, 32)  | 0       |
| max_pooling2d_3 (MaxPooling2D)              | (None, 46, 46, 32)  | 0       |
| dropout_4 (Dropout)                         | (None, 46, 46, 32)  | 0       |
| conv2d_8 (Conv2D)                           | (None, 44, 44, 64)  | 18432   |
| batch_normalization_7 (Batch Normalization) | (None, 44, 44, 64)  | 256     |
| activation_7 (Activation)                   | (None, 44, 44, 64)  | 0       |
| conv2d_9 (Conv2D)                           | (None, 42, 42, 64)  | 36864   |
| batch_normalization_8 (Batch Normalization) | (None, 42, 42, 64)  | 256     |
| activation_8 (Activation)                   | (None, 42, 42, 64)  | 0       |
| max_pooling2d_4 (MaxPooling2D)              | (None, 21, 21, 64)  | 0       |
| dropout_5 (Dropout)                         | (None, 21, 21, 64)  | 0       |
| conv2d_10 (Conv2D)                          | (None, 19, 19, 128) | 73728   |
| batch_normalization_9 (Batch Normalization) | (None, 19, 19, 128) | 512     |
| activation_9 (Activation)                   | (None, 19, 19, 128) | 0       |

|  |                     |         |
|--|---------------------|---------|
| conv2d_11 (Conv2D)                           | (None, 17, 17, 128) | 147456  |
| batch_normalization_10 (Batch Normalization) | (None, 17, 17, 128) | 512     |
| activation_10 (Activation)                   | (None, 17, 17, 128) | 0       |
| max_pooling2d_5 (MaxPooling2D)               | (None, 8, 8, 128)   | 0       |
| dropout_6 (Dropout)                          | (None, 8, 8, 128)   | 0       |
| flatten_1 (Flatten)                          | (None, 8192)        | 0       |
| dense_2 (Dense)                              | (None, 256)         | 2097152 |
| batch_normalization_11 (Batch Normalization) | (None, 256)         | 1024    |
| activation_11 (Activation)                   | (None, 256)         | 0       |
| dropout_7 (Dropout)                          | (None, 256)         | 0       |
| dense_3 (Dense)                              | (None, 1)           | 257     |

```

=====
Total params: 2386689 (9.10 MB)
Trainable params: 2385345 (9.10 MB)
Non-trainable params: 1344 (5.25 KB)
-----

```

```

[ ]: model2.compile(Adam(learning_rate=0.0001), loss='binary_crossentropy',
                    metrics=['accuracy'])

history2 = model2.fit(train_gen, steps_per_epoch=train_steps,
                      validation_data=val_gen,
                      validation_steps=val_steps,
                      epochs=20, verbose=1,
                      callbacks=callbacks_list)

```

```

Epoch 1/20
6189/6189 [=====] - 4178s 675ms/step - loss: 0.4447 -
accuracy: 0.7994 - val_loss: 0.8523 - val_accuracy: 0.6592 - lr: 1.0000e-04
Epoch 2/20
6189/6189 [=====] - 4524s 731ms/step - loss: 0.3323 -
accuracy: 0.8582 - val_loss: 0.6956 - val_accuracy: 0.7221 - lr: 1.0000e-04
Epoch 3/20
6189/6189 [=====] - ETA: 0s - loss: 0.2955 - accuracy:

```

0.8780

Epoch 3: ReduceLROnPlateau reducing learning rate to 9.999999747378752e-06.

6189/6189 [=====] - 4339s 701ms/step - loss: 0.2955 -

accuracy: 0.8780 - val\_loss: 0.7715 - val\_accuracy: 0.7295 - lr: 1.0000e-04

Epoch 4/20

6189/6189 [=====] - ETA: 0s - loss: 0.2656 - accuracy:

0.8923

Epoch 4: ReduceLROnPlateau reducing learning rate to 9.999999747378752e-07.

Restoring model weights from the end of the best epoch: 2.

6189/6189 [=====] - 4452s 719ms/step - loss: 0.2656 -

accuracy: 0.8923 - val\_loss: 1.0507 - val\_accuracy: 0.6895 - lr: 1.0000e-05

Epoch 4: early stopping



```
[ ]: acc2 = history2.history['accuracy']
      val_acc2 = history2.history['val_accuracy']
      loss2 = history2.history['loss']
      val_loss2 = history2.history['val_loss']

      epochs2 = range(1, len(acc2) + 1)

      plt.plot(epochs2, acc2, 'bo', label='Training Accuracy')
      plt.plot(epochs2, val_acc2, 'b', label='Validation Accuracy')
      plt.title('Training and validation Accuracy')
      plt.legend()
      plt.figure()

      plt.plot(epochs2, loss2, 'bo', label='Training loss')
      plt.plot(epochs2, val_loss2, 'b', label='Validation loss')
      plt.title('Training and validation loss')
      plt.legend()
      plt.figure()
```

[ ]: <Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>

### 0.1.3 Results and Analysis

While this is not the most encouraging results it is a marked improvement on the first model and gives around a 75% accuracy. As I was training this I also started running into hardware limitations that were causing training to take a long time. Since that is a limiting factor I will use this as my best result and resolve to use better hardware in future attempts of improvement.

### 0.1.4 Conclusion

Below the file for the Kaggle results are published and I achieve a mark of 75.66% while this is below average for this competition due to hardware issues during the training process I will have to accept the results. The fact that the validation accuracy is declining at the end of the training is also worrying.

```
[ ]: test_dir = 'data/histopathologic-cancer-detection/test'
      #os.mkdir(test_dir)

      # create test_images inside test_dir
      test_images = os.path.join(test_dir, 'test_images')
      os.mkdir(test_images)

[ ]: test_list = os.listdir('data/histopathologic-cancer-detection/test')

      for image in test_list:

          fname = image

          # source path to image
          src = os.path.join('data/histopathologic-cancer-detection/test', fname)
          # destination path to image
          dst = os.path.join(test_images, fname)
          # copy the image from the source to the destination
          #shutil.copyfile(src, dst)

[ ]: # Here we change the path to point to the test_images folder.

      test_gen = datagen.flow_from_directory('data/histopathologic-cancer-detection/
      ↪test',

                                             target_size=(img_size,img_size),
                                             batch_size=1,
                                             class_mode='categorical',
                                             shuffle=False)
```

Found 57458 images belonging to 1 classes.

```
[ ]: num_test_images = 57458
      predictions = model2.predict(test_gen, steps = num_test_images, verbose=1)
```

57458/57458 [=====] - 503s 9ms/step



```
[ ]: df_preds = pd.DataFrame(predictions, columns=[ 'has_tumor_tissue'])

print(df_preds.head())
```

```
      has_tumor_tissue
0          0.295501
1          0.547640
2          0.160517
3          0.013173
4          0.003677
```

```
[ ]: test_filenames = test_gen.filenames

# add the filenames to the dataframe
df_preds['file_names'] = test_filenames

df_preds.head()

def extract_id(x):

    # split into a list
    a = x.split('/')
    # split into a list
    b = a[1].split('.')
    extracted_id = b[0]

    return extracted_id

df_preds['id'] = df_preds['file_names'].apply(extract_id)

df_preds.head()

y_pred = df_preds['has_tumor_tissue']

# get the id column
image_id = df_preds['id']

submission = pd.DataFrame({'id':image_id,
                           'label':y_pred,
                           }).set_index('id')

submission.to_csv('preds.csv', columns=['label'])

submission.head()
```

```
[ ]:                                label
id
00006537328c33e284c973d7b39d340809f7271b 0.295501
0000ec92553fda4ce39889f9226ace43cae3364e 0.547640
00024a6dee61f12f7856b0fc6be20bc7a48ba3d2 0.160517
000253dfaa0be9d0d100283b22284ab2f6b643f6 0.013173
000270442cc15af719583a8172c87cd2bd9c7746 0.003677
```